



Big Data Analytics

– Chapter 6: Neural Networks –

Prof. Dr. Emmanuel Müller
Chair of Data Science and Data Engineering

Additional Literature for this Chapter


© Jiawei Han, Micheline Kamber, Jian Pei: Data Mining – Concepts and Techniques, 4th ed., Morgan Kaufmann Publishers, 2023.

Slides credit: <http://hanj.cs.illinois.edu/bk4/>

Content Overview

- Neural Networks
 - Basic concepts
 - Improve Training of Deep Learning Models
- Applications of Neural Networks
 - Convolutional Neural Networks
 - Recurrent Neural Networks
 - Graph Neural Networks
- Interpretable Machine Learning

Chapter 10. Deep Learning

- ❑ Basic Concepts 
- ❑ Improve Training of Deep Learning Models
- ❑ Convolutional Neural Networks
- ❑ Recurrent Neural Networks
- ❑ Graph Neural Networks
- ❑ Summary

Inspired by a biological neuron






Image credit:

<http://cs231n.github.io/neural-networks-1/>

How to model a single neuron?



Perceptron: Predecessor of a Neural Network



$$\hat{y} = f(\mathbf{W}^T \mathbf{x}) = f(\sum W_i x_i + b)$$

Activation Function

Adding non-linearity to the model

Weights

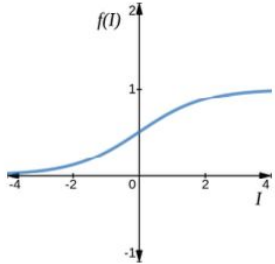
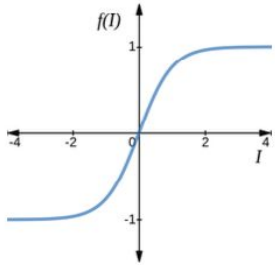
Bias

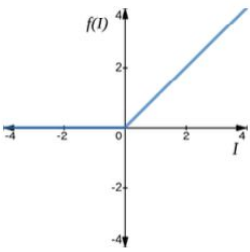
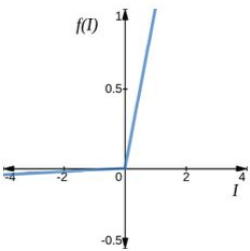
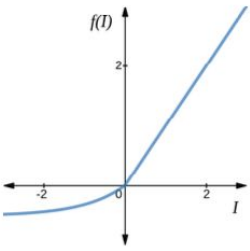
A measure of how easy it is to get the perceptron to output a 1

- ❑ Computes a weighted sum of inputs
- ❑ Invented in 1957 by Frank Rosenblatt. The original perceptron model does not have a non-linear activation function


Perceptron: Predecessor of a Neural Network

□ Examples of activation functions

Sigmoid	$\frac{1}{1+e^{-I}}$	
Tanh	$\frac{e^I - e^{-I}}{e^I + e^{-I}}$	


ReLU	$\begin{cases} 0, I \leq 0 \\ I, I > 0 \end{cases}$	
Leaky ReLU	$\begin{cases} 0.01 \times I, I < 0 \\ I, I \geq 0 \end{cases}$	
ELU	$\begin{cases} \alpha(e^I - 1), I \leq 0 \\ I, I > 0 \end{cases}$	

Multilayer Perceptron (MLP)




- ❑ Stacking multiple layers of perceptrons (adding hidden layers) makes a multilayer perceptron (MLP)
- ❑ MLP can engage in sophisticated decision making, where perceptrons fail
 - ❑ E.g. XOR problem
- ❑ Try it: <http://playground.tensorflow.org>

Feed-forward Neural Networks




Why multiple layers


- Automatic feature/representation learning
- Learn complicate (nonlinear) mapping function



(a) Input 4 tuples




(b) MLP




(c) Outputs of two hidden units

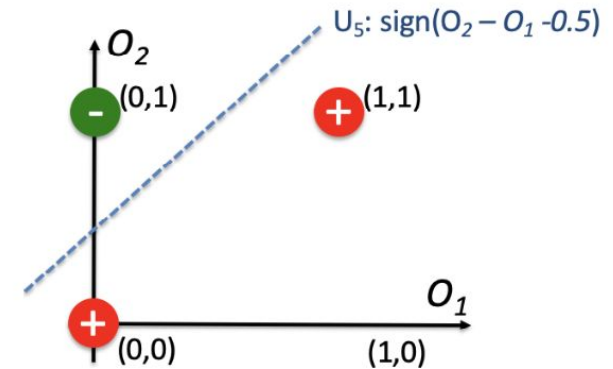
Feed-forward Neural Networks



(a) Input 4 tuples



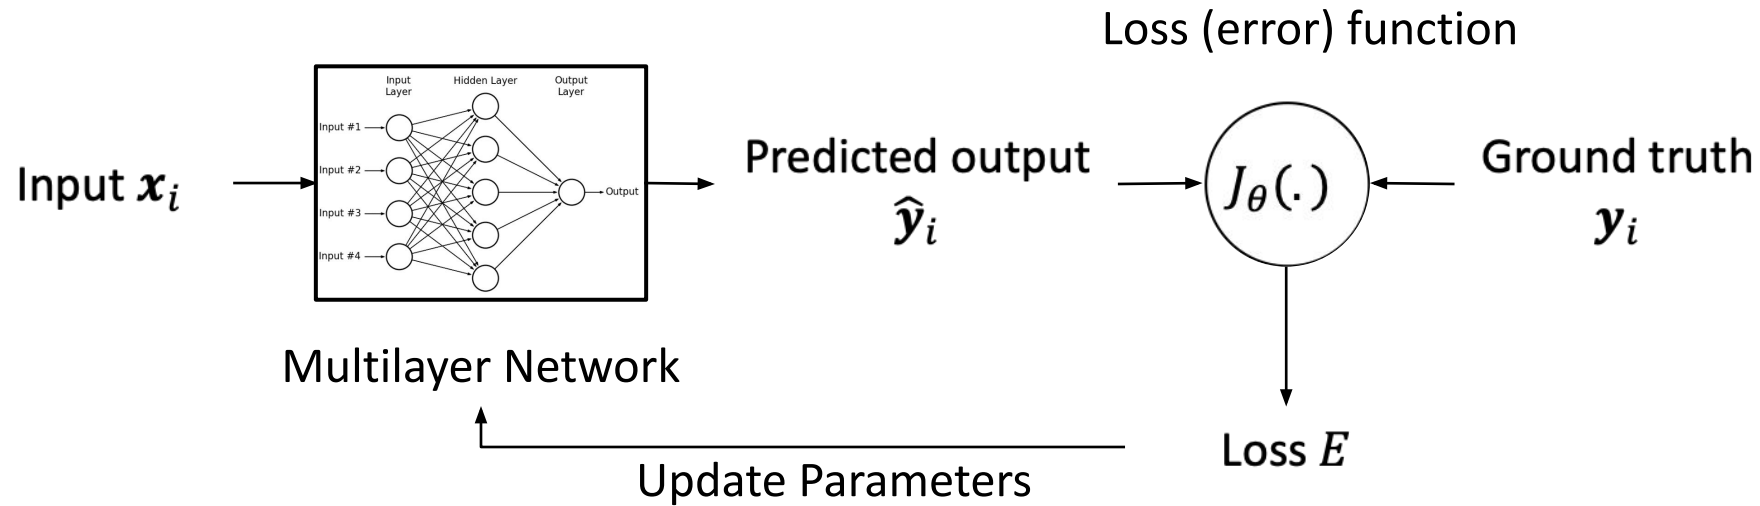
(b) MLP



(c) Outputs of two hidden units

x_1	x_2	$O_1 = \text{sign}(x_2 - x_1 - 0.5)$	$O_2 = \text{sign}(x_2 - x_1 + 0.5)$
0	0	0	+1
1	1	0	+1
1	0	0	0
0	1	1	1

Learning NN Parameters




□ Gradient Descent Algorithm

□ **Input:** Training sample x_i and its label y_i

1. **Feed Forward:** Get prediction $\hat{y}_i = \text{MLP}(x_i)$, and loss $E = J(\hat{y}_i, y_i)$
2. **Compute Gradient:** For each parameter θ_j (weights, bias), compute its gradient $\frac{\partial}{\partial \theta_j} J_\theta$
3. **Update Parameter:** $\theta_j = \theta_j - \eta \cdot \frac{\partial}{\partial \theta_j} J_\theta$

Empirical Explanation of Gradient Descent




$$\theta_j = \theta_j - \eta \cdot \frac{\partial}{\partial \theta_j} J_{\theta}$$

↓

η is the **learning rate**, which controls the 'step size' of the optimization


- ❑ Our objective is to minimize the loss function J , which is a function of the model parameters
- ❑ A gradient measures how much the output of a function changes if you change the inputs a little bit
- ❑ We update the parameters, based on their gradients, so that the loss function is going 'downhill'

Gradient Computation: Backpropagation




- The gradient of w_{ij} in the l th layer (corresponding to unit j in layer l , connected to unit i in layer $l-1$) is a function of
 - All 'error' terms from layer $l+1$ $\delta_k^{(l+1)}$ -- An auxiliary term for computation, not to be confused with gradients
 - Output from unit i in layer $l-1$ (input to unit j in layer l) -- Can be stored at the feed forward phase of computation

Gradient Computation: Backpropagation




- The 'error' terms $\delta_j^{(l)}$ is a function of
 - All $\delta_k^{(l+1)}$ in the layer $l+1$, if layer l is a hidden layer
 - The overall loss function value, if layer l is the output layer
 - We can compute the error at the output, and distribute backwards throughout the network's layers (backpropagation)
- Each forward pass and backpropagation over all training samples is called one **epoch**.

Gradient descent



(a) Gradient descent


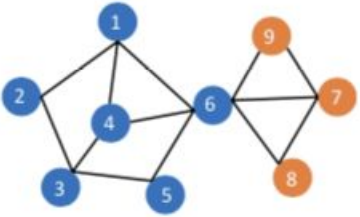
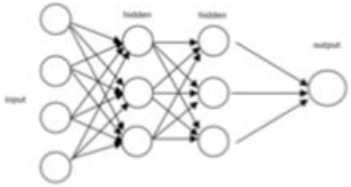
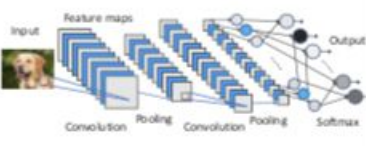
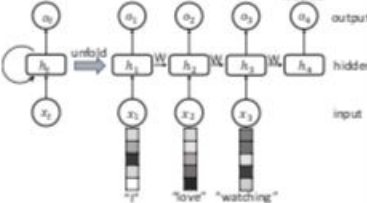
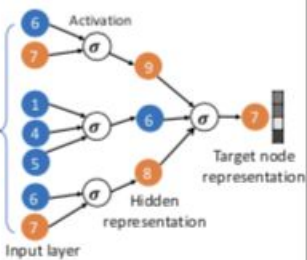


(b) Stochastic gradient descent


From Neural Networks to Deep Learning

- ❑ **Deep Learning** refers to training (deep) neural networks with many layers
 - ❑ More neurons, more layers
 - ❑ More complex ways to connect layers
- ❑ Deep Learning has some major advantages, making it popular
 - ❑ Tremendous improvement of performance in many tasks
 - ❑ Image recognition, natural language processing, AI game playing...
 - ❑ Requires no (or less) feature engineering, making end-to-end models possible
- ❑ Several factors lead to deep learning's success
 - ❑ Very large data sets
 - ❑ Massive amounts of computation power (GPU acceleration)
 - ❑ Advanced neural network structures and tricks
 - ❑ Convolutional neural networks, recurrent neural networks, graph convolutional networks
 - ❑ Dropout, ReLU, residual connection, ...

Overview of Typical Deep Learning Architectures

Data type	Multi-dimensional Features: credit rating, account balance $x = (4.5, 500, 3, 5)$ #deposits, #withdraws	Grid 	Sequence $x = \text{"I love watching movies."}$	Graph 
DL Architecture	Feed-forward Network 	CNN 	RNN 	GNN 

Chapter 10. Deep Learning

- ❑ Basic Concepts
- ❑ Improve Training of Deep Learning Models 
- ❑ Convolutional Neural Networks
- ❑ Recurrent Neural Networks
- ❑ Graph Neural Networks
- ❑ Summary

Techniques to Improve Deep Learning Training

- ❑ Key Challenges for Training Deep Learning Models
- ❑ Responsive Activation Functions
- ❑ Adaptive Learning Rate
- ❑ Dropout
- ❑ Pretraining
- ❑ Cross Entropy

Key Challenges

❑ Optimization Problem in Deep Learning

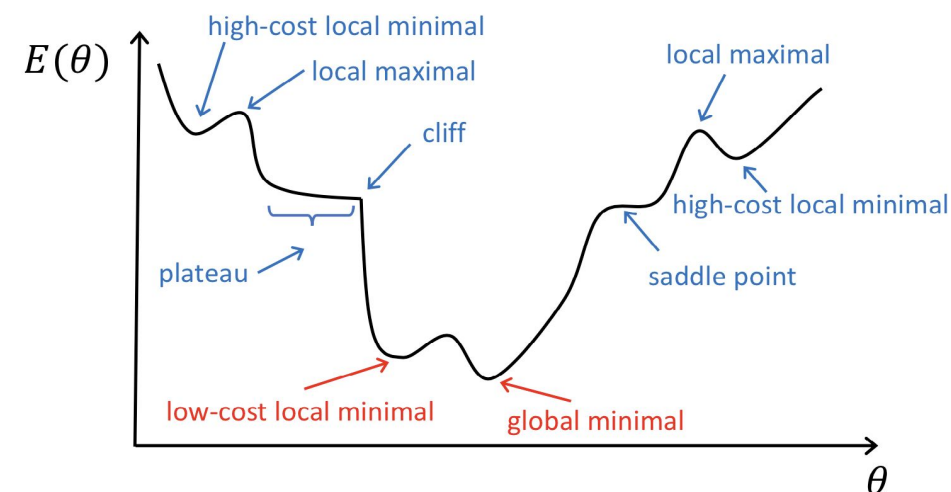
- ❑ Minimize the (approximated) training error E
- ❑ (Stochastic) gradient descent to find the model parameter

$$E(\boldsymbol{\theta}) = \frac{1}{m} \sum_{l=1}^m \text{Loss}(\hat{T}(X^l, \boldsymbol{\theta}), T^l)$$

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta \mathbf{g}_t$$

❑ Challenge 1: Optimization

- ❑ E is non-convex in general
- ❑ How to find a high-quality local optima



❑ Challenge 2: Generalization

- ❑ What we do: minimize (approximated) training error
- ❑ What we really want: minimize the generalization error
- ❑ How to mitigate over-fitting

Responsive Activation Function

❑ Saturation of Sigmoid Activation Function

- ❑ The output $\bar{O} = \sigma(I) = \frac{1}{1+e^{-I}} \in (0, 1)$
- ❑ The derivative $\frac{\partial O}{\partial I} = O(1 - O)$
- ❑ The error of an output unit $\delta_j = O_j(1 - O_j)(O_j - T_j)$

decaying

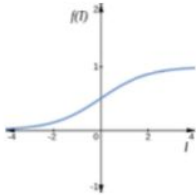
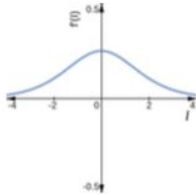
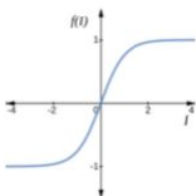
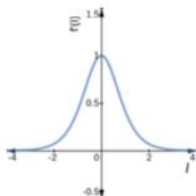
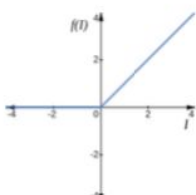
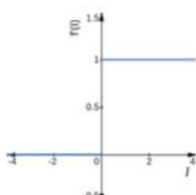
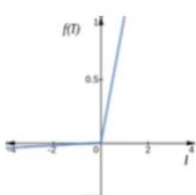
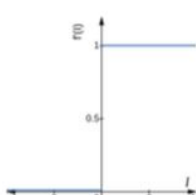
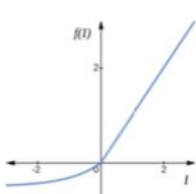

❑ Saturation of Sigmoid Activation Function

- ❑ If $O_j \approx 1$ or $O_j \approx 0$, both derivative and error will be close to 0
- ❑ Further exacerbated due to backpropagation -> gradient vanishing -> B.P. is stuck or takes long time to terminate


❑ A More Responsive Activation Function: Rectified Linear Unit (RELU)

- ❑ The output $O=f(I) = I$ if $I>0$, $O=0$ otherwise
- ❑ The gradient: $\frac{\partial O}{\partial I} = 1$ if $I > 0$ and $\frac{\partial O}{\partial I} = 0$
- ❑ The error: 0 if the unit is inactive ($I<0$), otherwise, aggregate all error terms from the units in the next higher layer the unit is connected to, **w/o decaying** -> avoid gradient vanishing

Activation Functions

Name	Definition ($f(I)$)	Plot	Derivative of $f(I)$	Plot
Sigmoid	$\frac{1}{1+e^{-I}}$		$f(I)(1 - f(I))$	
Tanh	$\frac{e^I - e^{-I}}{e^I + e^{-I}}$		$1 - f(I)^2$	
ReLU	$\begin{cases} 0, I \leq 0 \\ I, I > 0 \end{cases}$		$\begin{cases} 0, I \leq 0 \\ 1, I \geq 0 \end{cases}$	
Leaky ReLU	$\begin{cases} 0.01 \times I, I < 0 \\ I, I \geq 0 \end{cases}$		$\begin{cases} 0.01, I < 0 \\ 1, I \geq 0 \end{cases}$	
ELU	$\begin{cases} \alpha(e^I - 1), I \leq 0 \\ I, I > 0 \end{cases}$		$\begin{cases} \alpha e^I, I \leq 0 \\ 1, I > 0 \end{cases}$	

Adaptive Learning Rate




□ Stochastic gradient descent to find a local optima

- Default choice for η : a fixed, small positive constant $\theta_{t+1} = \theta_t - \eta g_t$
- Problems: slow progress, or jump over 'gradient cliff', or oscillation

□ Adaptive learning rate: let η change over epoch

- Strategy #1: $\eta_t = \frac{1}{t} \eta_0$
- Strategy #2: $\eta_t = (1 - \frac{t}{T})\eta_0 + \frac{t}{T}\eta_\infty$ if $t \leq T$ and $\eta_t = \eta_\infty$ otherwise



- Intuitions: smaller adjustment as algorithm progresses
- e.g., $\eta_0 = 0.9$; $\eta_\infty = 10^{-9}$
- Strategy #3 (AdaGrad): $\eta_t = \frac{1}{\rho + r_i} \eta_0$ $r_i = \sqrt{\sum_{k=1}^{t-1} g_{i,k}^2}$
 - Intuition: The magnitude of gradient g_t : indicator of the overall progress
 - e.g., $\rho = 10^{-8}$
- Strategy #4 (RMSProp): exponential decaying weighted sum of squared historical gradients

Dropout

□ The Purpose of Dropout: to Prevent Overfitting


□ How does it work?

- At each epoch (ρ dropout rate, e.g., $\rho = 0.5$)
 - randomly dropout some non-output units
 - Perform B.P. on the dropout network
- Scale the final model parameters


$$\theta^* \leftarrow \rho \cdot \theta^*$$

□ Why does Dropout Work?


- Dropout can be viewed as a regularization technique
 - Force the model to be more robust to noise, and to learn more generalizable features
- Dropout vs. Bagging
 - Bagging: Each base model is trained *independently* on a bootstrap sample
 - Dropout: the model parameters of the current dropout network are updated based on that of the previous dropout network(s)



(a) Original network




(b) Dropout U_3




(c) Dropout network

Pretraining


- ❑ **Pretraining:** the process of initializing the model in a suitable region
- ❑ **Greedy supervised pretraining**
 - ❑ pre-set the model parameters layer-by-layer in a greedy way
 - ❑ Start with simple model, add one additional layer at a time,
 - ❑ pretrain the parameters of the newly added layer, while fixing for other layers
 - ❑ Each time, equivalent to training a two-layered network
 - ❑ Followed up a fine—tuning process
- ❑ **Other Pretraining Strategies**
 - ❑ Unsupervised pretraining: based on auto-encoder
 - ❑ Hybrid strategy
- ❑ **Pretraining for transfer learning**




(a) original model to pretrain



(b) first iteration (1 hidden layer)



(c) second iteration (2 hidden layers)



(d) third iteration (3 hidden layers)

Cross Entropy

❑ Measure the disagreement between the actual (T) and predicted (O) target values

- ❑ For regression: mean-squared error
- ❑ For (binary) classification: cross-entropy


❑ Mean-squared error vs. Cross-entropy

	Mean-squared error	Cross entropy
Loss	$\frac{1}{2}(T - O)^2$	$-T \log O - (1 - T) \log (1 - O)$
Error δ	$O(1 - O)(O - T)$	$O - T$

❑ Cross-entropy for Multiclass Problem

- ❑ Actual target $T = (T_1, T_2, \dots, T_C)$
- ❑ Predicted output $O = (O_1, O_2, \dots, O_C)$

$$\text{Loss}(T, O) = - \sum_{j=1}^C T_j \log O_j$$




Assume positive example ($T=1$)

Example: House Price Prediction

Task: predicting house prices with an MLP (optimize Mean-squared error)

- Input layer (3 neurons): each represents a feature in the input data
- Hidden layer (5 neurons): learned weights capture relationships within the data
- Output layer (1 neuron): produces continuous output prediction based on learned patterns

$X_1 =$ Area	$X_2 =$ # Bedrooms	$X_3 =$ Air-conditioning	$y =$ House Price
80	1	No (0)	250.000
200	3	Yes (1)	600.000
...
180	2	Yes (1)	?



Example: House Price Classification

Transformation into binary classification problem:

- Discretize house prices into C1 = Affordable, C2 = Too Expensive (e.g., use fixed threshold)
- Adjust MLP architecture:
 - One output neuron for each class, providing raw class scores (logits)
 - Softmax function turns logits into class probabilities
 - Optimize Cross-entropy w.r.t. class probabilities

